

Exotic Options

In this lecture...

- the names and contract details for many basic types of exotic options
- how to classify exotic options according to important features
- how to think about derivatives in a way that makes it easy to compare and contrast different contracts
- pricing exotics using Monte Carlo simulation
- pricing exotics via partial differential equations and then finite-difference methods

By the end of this lecture you will be able to

- characterize most exotic contracts according to a list of important features
- price exotics using Monte Carlo simulations
- interpret the pricing of many exotics in terms of partial differential equations

Introduction

Exotic contracts are traded **over the counter (OTC)**, meaning that they are designed by the relevant counterparties and are not available as exchange-traded contracts.

Exotic options include contracts with features making them more complex to price and to hedge than vanillas.

Often one takes volatilities 'implied' by the market prices of vanillas and put them into the pricing model for exotics.

Note: Any 'new theories' that we build up will be done in an equity framework for simplicity. But the ideas apply to all markets (and we'll see term sheets from various markets later.)

An example of an exotic option

For Discussion Purposes Only

Five-year Minimum Coupon Cliquet on SPX Index

| | |
|---------------------------------------|--|
| Option Buyer | XXXX |
| Option Seller | YYYY |
| Notional Amount | \$ 25MM |
| Trade Date | 20 December 2000 |
| Start Date | 31 January 2001 |
| Maturity Date | Start Date + Five years |
| Option Seller Pays at Maturity | Notional * $\max \left(\sum_{i=1}^5 \max \left(0, \min \left(\text{Cap}, \frac{S_i - S_{i-1}}{S_{i-1}} \right) \right), 16.25\% \right)$ |
| Index | SPX Index |
| Cap | 8% |
| Option Premium | ? |
| Index Levels | S_i = Closing Level of Index on Start Date + i years |

This indicative termsheet is neither an offer to buy or sell securities or an OTC derivative product which includes options, swaps, forwards and structured notes having similar features to OTC derivative transactions, nor a solicitation to buy or sell securities or an OTC derivative product. The proposal contained in the foregoing is not a complete description of the terms of a particular transaction and is subject to change without limitation.

Important features to look out for

- Time dependence
- Cashflows
- Path dependence
- Dimensionality
- Order
- Embedded decisions

1. Time dependence

Here we are concerned with time dependence in the option contract.

For example, discrete cashflows necessarily involve time dependence.

Another example, early exercise might only be permitted on certain dates or during certain periods. This intermittent early exercise is a characteristic of **Bermudan options**.

Similarly, the position of the barrier in a knock-out option may change with time. Every month it may be reset at a higher level than the month before.

- These contracts are referred to as **time inhomogeneous**.

When there is time dependence in a contract we might expect

- jumps in option values and/or the greeks
- to have to worry about time discretization in numerical schemes

2. Cashflows


Imagine a contract that pays the holder an amount q at time t_i . The contract could be a bond and the payment a coupon.

If we use $V(t)$ to denote the contract value and t_i^- and t_i^+ to denote just before and just after the cashflow date then simple arbitrage considerations lead to

- $$V(t_i^-) = V(t_i^+) + q.$$

This is a **jump condition**.

The value of the contract jumps by the amount of the cashflow. The behavior of the contract value across the payment date is shown in the figure.



A discrete cashflow and its effect on a contract value.

If the contract is contingent on an underlying variable so that we have $V(S, t)$ then we can accommodate cashflows that depend on the level of the asset S i.e. we could have $q(S)$.

Certificate in Quantitative Finance

That's an example of a **discrete cashflow**.

Some contracts specify **continuous cashflows**. There may be a payment every day.

When there are cashflows we expect

- option values to jump
- the greeks to jump

3. Path dependence

Path-dependent contracts have payoffs, and therefore values, that depend on the history of the asset price path.

An asset starts at A and ends at Z at expiration. If the contract is path dependent the route taken from A to Z matters. If it is not path dependent then the route does not matter.

Path dependence comes in two main forms:

- Strong path dependence
- Weak path dependence

Strong path dependence

Of particular interest, mathematical and practical, are the **strongly path-dependent contracts**.

These have payoffs that depend on some property of the asset price path in addition to the value of the underlying at the present moment in time; in the equity option language, we cannot write the value as $V(S, t)$.

- The contract value is a function of at least one more independent variable.


Example:

The Asian option has a payoff that depends on the average value of the underlying asset from inception to expiry. We must keep track of more information about the asset price path than simply its present position.

The extra information that we need is contained in the 'running average.' This is the average of the asset price from inception until the present, when we are valuing the option.



Path dependency also comes in **discrete** and **continuous** varieties depending on whether the path-dependent quantity is **sampled** discretely or continuously.



When there is strong path dependence in a contract we might expect

- to have to solve in higher dimensions


(We have to keep track of a new **state variable** such as the average to date.)

Weak path dependence

- Options whose value depends on the asset history, but can still be written as $V(S, t)$ are said to be **weakly path dependent**.

One of the most common reasons for weak path dependence in a contract is a **barrier**. Barrier (or knock-in, or knock-out) options are triggered by the action of the underlying hitting a prescribed value at some time before expiry.

For example, as long as the asset remains below 150, the contract will have a call payoff at expiry. However, should the asset reach this level before expiry then the option becomes worthless; the option has 'knocked out.'



Two paths having the same value at expiry but with completely different payoffs.

Certificate in Quantitative Finance

- Weak path dependency does not add any extra dimensions.

(So a barrier option still only has two dimensions, S and t .)

4. Dimensionality

Dimensionality refers to the number of underlying independent variables.

- The vanilla option has two independent variables, S and t , and is thus two dimensional.
- The weakly path-dependent contracts have the same number of dimensions as their non-path-dependent cousins, i.e. a barrier call option has the same two dimensions as a vanilla call.

But some contracts require us to go in to extra dimensions!

There are two distinct reasons why we need more dimensions. . .

- More sources of randomness
- Strong path dependence

More dimensions caused by more sources of randomness

We will get higher dimensions if we have more sources of randomness

- If we have an option on **10** underlyings ('best of' for example) we will have **11** dimensions (S_1, S_2, \dots, S_{10} and t)

But we will also get more dimensions if we have other types of randomness, such as volatility.

- If we have an option on **10** underlyings and we use a stochastic volatility model for each asset we will have **21** dimensions (S_1, S_2, \dots, S_{10} and t , and also $\sigma_1, \sigma_2, \dots, \sigma_{10}$)

Each new dimensions introduces extra 'diffusion' terms. (What does this mean for the governing PDE?)

How many parameters will there be?

(GIGO?)

More dimensions caused by strong path dependency

We will get higher dimensions if we have an option that is strongly path dependent.

- If we have an option that pays off the maximum of the average stock price we will have **4** dimensions (S and t , but also a state variable for the average and another for the maximum of the average!)

We'll see the theory of this later, but the effect on the governing PDE is to sometimes add new terms that are not diffusive! (But no new parameters!)

When the problem is of high dimensions we might expect

- to have restrictions on the kind of numerical solution we employ. The higher the number of dimensions, the more likely we are to want to use Monte Carlo simulations.

5. The order of an option

The basic, vanilla options are of first order. Their payoffs depend only on the underlying asset, the quantity that we are *directly* modeling. Other, path-dependent, contracts can still be of first order if the payoff only depends only on properties of the asset price path.

- **Higher order** refers to options whose payoff, and hence value, is contingent on the value of *another* option.

The obvious second-order options are compound options, for example, a call option giving the holder the right to buy a put option. The compound option expires at some date T_1 and the option on which it is contingent, expires at a later time T_2 . Technically speaking, such an option is weakly path dependent.

From a practical point of view, the compound option raises some important modeling issues.

- The payoff for the compound option depends on the *market* value of the underlying option, and not on the theoretical price.

If you hold a compound option, and want to exercise the first option then you must take possession of the underlying option. High order option values are very sensitive to the basic pricing model and should be handled with care.

When options are of higher order we might expect

- to have to worry more about the accuracy of the model
- to have to solve a multi-layer problem, first for the 'underlying option' and then for the exotic

6. Decisions

Holding an American option you are faced with the decision whether and when to exercise your rights.

The American option is the most common contract that contains within it a decision feature. Other contracts require more subtle and interesting decisions to be made.

For example, the conversion of convertible bonds is mathematically identical to the early exercise of an American option.

Example: The passport option is an option on a trading account.

You buy and sell some asset, if you are in profit on the expiry of the option you keep the money, if you have made a loss it is written off.

The decisions to be made here are when to buy, sell or hold, and how much to buy, sell or hold.

Embedded decisions are common features.

- The key point about any decision, including early exercise, is that the holder of this valuable right should ideally act *optimally*, i.e. they must decide *when* to exercise the option, convert the bond etc.

What is meant by 'optimal'?

This is best explained by an example.

Example: You want to sell an American option with one week until expiration.

You know how to delta hedge to eliminate market risk, sensitivity to the direction of the underlying asset.

But you are still exposed to the exercise strategy of the option holder.


How can you price the option? What must you assume about the exercise strategy of the holder?

You price assuming that the holder **exercises in such a way as to give the option its highest value.**

- You sell the option for the highest attainable value over all possible decision strategies.

That sounds complicated. But we have already seen some of the mathematics for this idea.

The end result is that the option value must always be above the payoff.

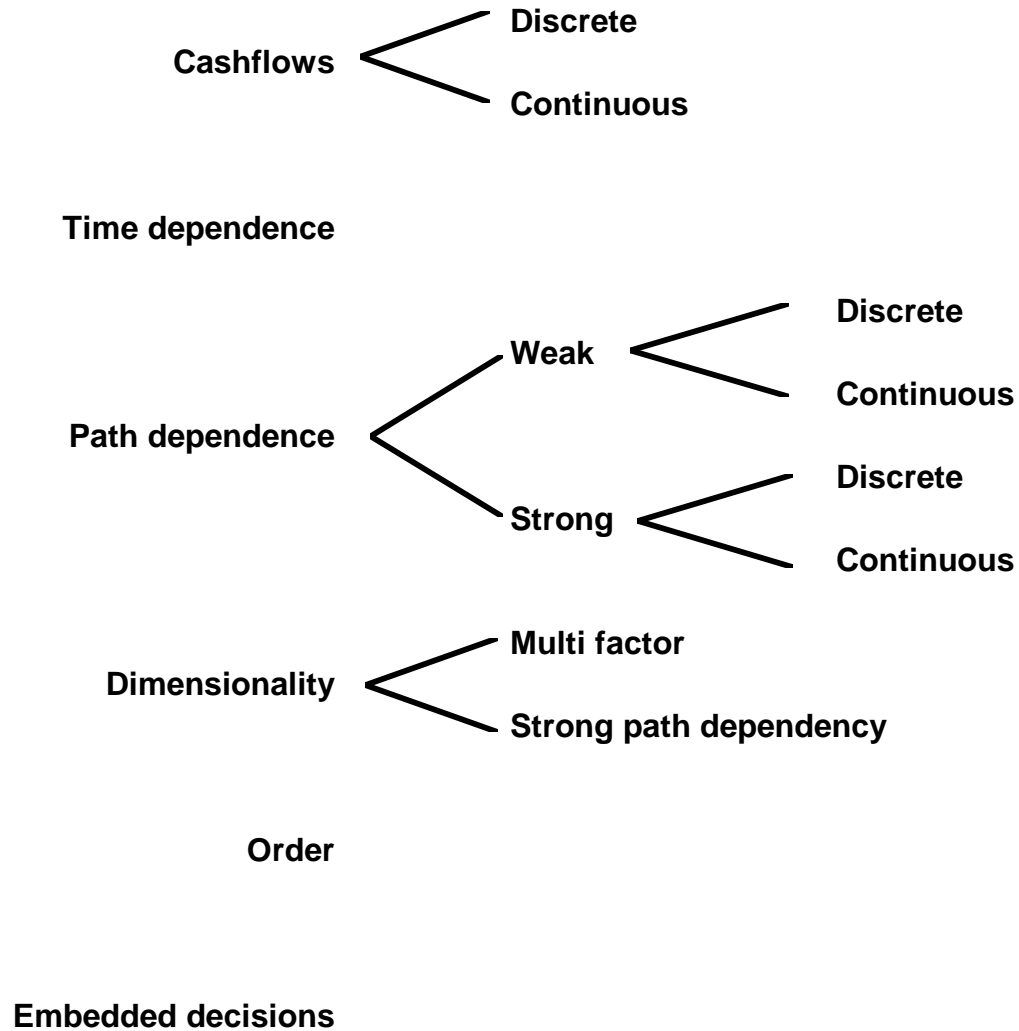


And if the issuer has to make some decisions?

When options have embedded decisions we might expect

- to have to model the optimal decision, from either the holder's or issuer's perspective (or both)
- to have restrictions on the numerical solution. Monte Carlo is not a good method for pricing options in which there are embedded decisions.
- to have a \geq or \leq term in the model/code

Schematic diagram of exotic option classification system:



Pricing methodologies

Now let's look at the (numerical) pricing of exotic options.

Two main methods:

- Pricing via simulations, Monte Carlo
- Formulating the pricing problem in terms of partial differential equations, for solving by finite-difference methods

Pricing via expectations, Monte Carlo simulation

We can value options in the Black–Scholes world by taking the present value of the expected payoff under a risk-neutral random walk.

Simply simulate the random walk

$$dS = rS dt + \sigma S dX$$

for many paths, calculate the payoff for each path—and this means calculating the value of the path-dependent quantity which is usually very simple to do—take the average payoff over all the paths and then take the present value of that average.

That is the option's fair value.

Certificate in Quantitative Finance

When and when not to use MC

This is a very general and powerful technique.

When to use MC:

- Good when there are a large number of dimensions
- Useful for path-dependent contracts (even if low dimensions!) for which a partial differential equation approach is tedious to set up
- Some models (e.g. HJM) are built for MC, not easy (or impossible) to write as PDE

When not to use MC:

- The main disadvantage is that it is hard to value options with embedded decisions using MC simulation

Partial differential equations and finite differences

To be able to turn the valuation of a derivatives contract into the solution of a partial differential equation is a big step forward.

- The partial differential equation approach is one of the best ways to price a contract because of its flexibility and because of the large body of knowledge that has grown up around the fast and accurate numerical solution of these problems
- But there is effort involved in setting up the PDE for numerical solution. (In contrast, Monte Carlo can be used 'straight out of the bag')

Let's look at setting up the PDE approach for two examples, a **barrier** option and an **Asian** option.

Both of these can be priced via Monte Carlo but finite-difference solution of the PDEs will be faster.

- Is it worth the effort? Sometimes you might do initial pricing via MC (just to get a 'number') and then you'll spend a bit of time coding up finite differences before it goes into the bank's 'system.'

After we've looked at these two problems we'll do a general theory of path-dependent options.

Barrier options


- **Barrier options** have a payoff that is contingent on the underlying asset reaching some specified level before expiry.

The critical level is called the barrier, there may be more than one.

Barrier options come in two main varieties, the 'in' barrier option (or **knock-in**) and the 'out' barrier option (or **knock-out**). The former only have a payoff if the barrier level is reached before expiry and the latter only have a payoff if the barrier is *not* reached before expiry.


These contracts are weakly path dependent.

Example: An up-and-out call option. This has a call payoff at expiration unless the barrier has been triggered some time before expiration.



This is easily solved by Monte Carlo simulation or by finite-difference methods.

The latter is much preferable. And because the barrier option is weakly path dependent the relevant PDE is exactly the classical Black–Scholes equation! We just have to figure out **initial** and **boundary conditions**.



Asian options

- **Asian options** have a payoff that depends on the average value of the underlying asset over some period before expiry.

They are strongly path dependent. Their value prior to expiry depends on the path taken.

The average used in the calculation of the option's payoff can be defined in many different ways.

It can be an **arithmetic average** or a **geometric average**, for example.

The data could be **continuously sampled**, so that every realized asset price over the given period is used. More commonly, for practical and legal reasons, the data is usually **sampled discretely**.

How is the continuously sampled arithmetic average, for example, defined mathematically?



The final payoff is a function of

$$A = \frac{1}{T} \int_0^T S(\tau) d\tau.$$

(The averaging started at time $t = 0$.)

But the running average, and hence our new **state variable** is

$$A = \frac{1}{t} \int_0^t S(\tau) d\tau.$$

We need a theory for options with payoff depending on integrals.

A theory for strong path dependence

We will now look at

- pricing many strongly path-dependent contracts in the Black–Scholes partial differential equation framework
- how to handle both continuously sampled and discretely sampled paths
- jump conditions for differential equations

We will now see how to generalize the Black–Scholes analysis, delta hedging and no arbitrage, to the pricing of many more derivative contracts, specifically contracts that are strongly path dependent.

Path-dependent quantities represented by an integral

We start by assuming that the underlying asset follows the log-normal random walk

$$dS = \mu S dt + \sigma S dX.$$

Imagine a contract that pays off at expiry, T , an amount that is a function of the path taken by the asset between time zero and expiry.

- Let us suppose that this path-dependent quantity can be represented by an integral of some function of the asset over the period zero to T :

$$I(T) = \int_0^T f(S, \tau) d\tau.$$

This is not such a strong assumption, many of the path-dependent quantities in exotic derivative contracts, such as averages, can be written in this form with a suitable choice of $f(S, t)$.

You might think that we need to model and remember S at every single moment between now and expiration.

This may look like a problem with an infinite number of variables.

It is much easier than this.

For the basic Asian option it turns out, as we shall see, that not all of the past of the asset matters, only one ‘functional’ of it.

Prior to expiry we have information about the possible final value of S (at time T) in the present value of S (at time t).

For example, the higher S is today, the higher it will probably end up at expiry.

Similarly, we have information about the possible final value of I in the value of the integral to date:

$$I(t) = \int_0^t f(S, \tau) d\tau. \quad (1)$$

As we get closer to expiry, so we become more confident about the final value of I .

- One can imagine that the value of the option is therefore not only a function of S and t , but also a function of I ; I will be our new independent variable, called a **state variable**.

We see in the next section how this observation leads to a pricing equation.

In anticipation of an argument that will use Itô's lemma, we need to know the stochastic differential equation satisfied by I .

This could not be simpler.

- Incrementing t by dt in (1) we find that

$$dI = f(S, t) dt. \quad (2)$$

Observe that I is a smooth function (except at discontinuities of f) and from (2) we can see that its stochastic differential equation does not contain any stochastic terms.

Continuous sampling: The pricing equation

We will derive the pricing partial differential equation for a contract that pays some function of our new variable I .

- The value of the contract is now a function of the three variables, $V(S, I, t)$.

Set up a portfolio containing one of the path-dependent option and short a number Δ of the underlying asset:

$$\Pi = V(S, I, t) - \Delta S.$$

The change in the value of this portfolio is given by

$$d\Pi = \left(\frac{\partial V}{\partial t} + \frac{1}{2} \sigma^2 S^2 \frac{\partial^2 V}{\partial S^2} \right) dt + \frac{\partial V}{\partial I} dI + \left(\frac{\partial V}{\partial S} - \Delta \right) dS.$$

- Choosing

$$\Delta = \frac{\partial V}{\partial S}$$

to hedge the risk, and using (2), we find that

$$d\Pi = \left(\frac{\partial V}{\partial t} + \frac{1}{2}\sigma^2 S^2 \frac{\partial^2 V}{\partial S^2} + f(S, t) \frac{\partial V}{\partial I} \right) dt.$$

This change is risk free, and thus earns the risk-free rate of interest r , leading to the pricing equation...

$$\frac{\partial V}{\partial t} + \frac{1}{2}\sigma^2 S^2 \frac{\partial^2 V}{\partial S^2} + f(S, t) \frac{\partial V}{\partial I} + rS \frac{\partial V}{\partial S} - rV = 0$$

This is to be solved subject to

$$V(S, I, T) = \text{Payoff}(S, I).$$

This completes the formulation of the valuation problem.

Example:

Continuing with the arithmetic Asian example, we have

$$I = \int_0^t S d\tau,$$

so that the equation to be solved is

$$\frac{\partial V}{\partial t} + \frac{1}{2}\sigma^2 S^2 \frac{\partial^2 V}{\partial S^2} + S \frac{\partial V}{\partial I} + rS \frac{\partial V}{\partial S} - rV = 0.$$

Similarity reductions

As long as the stochastic differential equation for the path-dependent quantity only contains references to S , t and the path-dependent quantity itself then the value of the option depends on three variables.

Unless we are very lucky, the value of the option must be calculated numerically.

- Some options have a particular structure that permits a reduction in the dimensionality of the problem by use of a similarity variable.

The dimensionality of the continuously sampled arithmetic average strike option can be reduced from three to two.

The payoff for the continuously sampled arithmetic average strike call option is

$$\max \left(S - \frac{1}{T} \int_0^T S(\tau) d\tau, 0 \right).$$

This can be written as

$$I \max \left(R - \frac{1}{T}, 0 \right)$$

where

$$I = \int_0^t S(\tau) d\tau$$

and

$$R = \frac{S}{\int_0^t S(\tau) d\tau}.$$

In view of the form of the payoff function, it seems plausible that the option value takes the form

$$V(S, I, t) = IW(R, t), \quad \text{with} \quad R = \frac{S}{I}.$$

We find that W satisfies

$$\frac{\partial W}{\partial t} + \frac{1}{2}\sigma^2 R^2 \frac{\partial^2 W}{\partial R^2} + R(r - R) \frac{\partial W}{\partial R} - (r - R)W = 0$$

with final condition

$$W(R, T) = \max\left(R - \frac{1}{T}, 0\right).$$

Path-dependent quantities represented by an updating rule

For practical and legal reasons path-dependent quantities are never measured continuously.

There is minimum time step between sampling of the path-dependent quantity.

- From a practical viewpoint it is difficult to incorporate every single traded price into an average, for example. Data can be unreliable and the exact time of a trade may not be known accurately.

If the time between samples is small we can confidently use a continuous-sampling model, the error will be small.

If the time between samples is long, or the time to expiry itself is short we must build this into our model. This is the goal of this section.

When path-dependent quantities are sampled discretely we have summations instead of integrals.

Example: The discretely sampled Asian option

We saw how to use the continuous running integral in the valuation of Asian options.

But what if that integral is replaced by a discrete sum?

In practice, the payoff for an Asian option depends on

$$A_M = \frac{I_M}{M} = \frac{1}{M} \sum_{k=1}^M S(t_k), \quad (3)$$

where M is the total number of sampling dates.

This is the discretely sampled average.

This is simply a discrete version of our earlier continuous integral.

Example: The payoff for a discretely sampled arithmetic average strike put is then

$$\max(A_M - S, 0).$$

Must we remember every single $S(t_k)$ to price the option? That would be an $M + 2$ -dimensional problem!

Recall that when valuing the continuously sampled Asian option we only had to remember the value of a single new quantity, the average to date.

Is the same true with the discretely sampled Asian?

Can we write the expression for the running discretely sampled average in a form that does not require us to remember every single $S(t_k)$?

Yes.

$$A_i = \frac{1}{i} \sum_{k=1}^i S(t_k)$$

So that

$$A_1 = S(t_1),$$

$$A_2 = \frac{S(t_1) + S(t_2)}{2} = \frac{1}{2}A_1 + \frac{1}{2}S(t_2),$$

$$A_3 = \frac{S(t_1) + S(t_2) + S(t_3)}{3} = \frac{2}{3}A_2 + \frac{1}{3}S(t_3), \dots$$

- This can be expressed as an **updating rule**

$$A_i = \frac{1}{i}S(t_i) + \frac{i-1}{i}A_{i-1}.$$

Generalization

An **updating rule** is an algorithm for defining the path-dependent quantity in terms of the current 'state of the world.'

- The path-dependent quantity is measured on the **sampling dates** t_i , and takes the value I_i for $t_i \leq t < t_{i+1}$.
- At the sampling date t_i the quantity I_{i-1} is updated according to a rule such as

$$I_i = F(S(t_i), I_{i-1}, i).$$

Note how, in this simplest example, the new value of I is determined by only the old value of I and the value of the underlying on the sampling date, and the sampling date.

Another example: the Lookback option

We will see how to use this for pricing in the next section. But first, another example.

The lookback option has a payoff that depends on the maximum or minimum of the realized asset price.



If the payoff depends on the maximum sampled at times t_i then we have

$$I_1 = S(t_1), \quad I_2 = \max(S(t_2), I_1), \quad I_3 = \max(S(t_3), I_2) \cdots$$

- The updating rule is therefore simply

$$I_i = \max(S(t_i), I_{i-1}).$$

How do we use these updating rules in the pricing of derivatives?

Discrete sampling: The pricing equation

- We anticipate that the option value will be a function of three variables, $V(S, I, t)$.

The first step in the derivation is the observation that the stochastic differential equation for I is degenerate:

$$dI = 0.$$

This is because the variable I can only change at the discrete set of dates t_i . This is true if $t \neq t_i$ for any i .

- So provided we are not *on* a sampling date the quantity I is constant, the stochastic differential equation for I reflects this, and the pricing equation is simply the basic Black–Scholes equation:

$$\frac{\partial V}{\partial t} + \frac{1}{2}\sigma^2 S^2 \frac{\partial^2 V}{\partial S^2} + rS \frac{\partial V}{\partial S} - rV = 0.$$

How does the equation know about the path dependency?

- Across a sampling date the option value is continuous.

As we get closer and closer to the sampling date we become more and more sure about the value that I will take according to the updating rule.

Since the outcome on the sampling date is known and since *no money changes hands* there cannot be any jump in the value of the option.

This is a simple application of the no arbitrage principle.

We introduce the notation t_i^- to mean the time infinitesimally before the sampling date t_i and t_i^+ to mean infinitesimally just after the sampling date.

Continuity of the option value is represented mathematically by


$$V(S, I_{i-1}, t_i^-) = V(S, I_i, t_i^+).$$

- In terms of the updating rule, we have

$$V(S, I, t_i^-) = V(S, F(S(t_i), I, i), t_i^+)$$

This is called a **jump condition**.

We call this a jump condition even though there is no jump in the option value!



Examples:

- To price an arithmetic Asian option with the average sampled at times t_i solve the Black–Scholes equation for $V(S, A, t)$ with

$$V(S, A, t_i^-) = V\left(S, \frac{i-1}{i}A + \frac{1}{i}S, t_i^+\right).$$

- To price a lookback depending on the maximum sampled at times t_i solve the Black–Scholes equation for $V(S, M, t)$ with


$$V(S, M, t_i^-) = V\left(S, \max(S, M), t_i^+\right).$$


Let's see how this is applied.

The top right-hand plot is the S, M plane just after the sample of the maximum has been taken.

Because the sample has just been taken the region $S > M$ cannot be reached, it is the region labeled 'Unattainable.'

This means that the option value at time t_i^- for $S < M$ is the same as the t_i^+ value. However, for $S > M$ the option value comes from the $S = M$ line at time t_i^+ for the same S value.





The algorithm for discrete sampling

The path-dependent quantity, I , is updated discretely and so the partial differential equation for the option value between sampling dates is the Black–Scholes equation.

The algorithm for valuing an option on a discretely sampled quantity is as follows.

- Working backwards from expiry, solve

$$\frac{\partial V}{\partial t} + \frac{1}{2}\sigma^2 S^2 \frac{\partial^2 V}{\partial S^2} + rS \frac{\partial V}{\partial S} - rV = 0$$

between sampling dates. Stop when you get to the time step on which the sampling takes place.

- Then apply the appropriate jump condition across the current sampling date to deduce the option value immediately before the present sampling date using the calculated value of the option just after. Use this as your final condition for further time stepping of the Black–Scholes equation.
- Repeat this process to arrive at the current value of the option.

When and when not to use PDEs/FD

PDEs and finite differences take more effort to set up than MC but the reward can be in speed, etc.

When to use PDEs/FD:

- These techniques are best for low dimensions
- They handle embedded decisions exceptionally well

When not to use PDEs/FD:

- When there are high dimensions finite differences will struggle
- Some path dependency can be difficult to model

Now let's look at some term sheets!

(And please take a look at the “Extras” on exotics.)

Summary

Please take away the following important ideas

- Exotics can be easily characterized by important features
- As long as they don't have embedded decisions they can usually be priced by Monte Carlo simulations
- Many exotics can be understood via partial differential equations, which may help in their efficient pricing